

# DEEP GENERATIVE STOCHASTIC NETWORKS FOR SEQUENCE PREDICTION

Markus B. Beissinger

A THESIS

in

Computer and Information Science

Presented to the Faculties of the University of Pennsylvania in Partial Fulfillment  
of the Requirements for the Degree of Master of Science in Engineering

2014

---

Lyle H. Ungar  
Supervisor of Thesis

---

Val Tannen  
Graduate Group Chairperson

# Acknowledgements

I would like to thank Lyle Ungar for his invaluable insights and methods of thinking throughout advising this thesis. I would also like to thank Mark Liberman (CIS, Linguistics, University of Pennsylvania) and Mitch Marcus (CIS, University of Pennsylvania) for their work on the thesis committee. Finally, I would like to thank Li Yao (Computer Science and Operations Research, Université de Montréal) from Yoshua Bengio's group for the Theano framework and help with starter code.

# Abstract

This thesis presents three methods of using Generative Stochastic Networks (GSN) for sequence prediction. GSNs are a recent probabilistic generalization of denoising auto-encoders that learn unsupervised deep hierarchical representations of complex input data, while being trainable by backpropagation. Each method presented relies on learning the transition operator on a Markov chain of the input data over time. The first method (Model 1) views GSNs as learning complex representations of the individual input data using latent state variables  $H$ , so that a simple regression step  $H \rightarrow H$  can encode the next set of latent state variables describing the next input data in the sequence, learning  $P(X_{t+1}|X_t, H_t)$ . This method is similar to Expectation Maximization (EM), where first a GSN is trained on all input data, and then a regression is trained on the sequenced latent states  $H$ , which is repeated until training converges. The second method (Model 2) is an online training method that views GSNs as recurrent networks that encode sequences of input over time. This means the networks learn features of the input data that best predict the next expected data in the sequence. The third method (Model 3) uses a hybrid approach of combining GSN layers with a recurrent layer. This allows the GSN to learn useful features of the input data while the recurrent network learns the sequence representation over these features. Experimental results for these three methods are presented on artificially sequenced handwritten digit MNIST data (digits sequenced 0-9 repeating) and data generated from the models.

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Related Work</b>	<b>3</b>
<b>3</b>	<b>Background: Deep Architectures</b>	<b>4</b>
3.1	Representation Learning . . . . .	4
3.2	Interpretations of Deep Architectures . . . . .	5
3.2.1	Probabilistic models: restricted boltzmann machine (RBM) .	7
3.2.2	Direct encoding models: auto-encoder . . . . .	9
3.3	Denoising Auto-encoders . . . . .	11
<b>4</b>	<b>General Methodology: Deep Generative Stochastic Networks (GSN)</b>	<b>13</b>
4.1	Generalizing denoising auto-encoders . . . . .	13
4.1.1	Easing restrictive conditions on the denoising auto-encoder .	14
4.1.2	Generalizing to GSN . . . . .	14
4.2	Extension to recurrent deep GSN . . . . .	16
<b>5</b>	<b>Model 1: Expectation Maximization method for training a recurrent deep GSN</b>	<b>17</b>
5.1	Recurrent nature of deep GSNs . . . . .	17
5.2	Algorithm . . . . .	17
5.3	Experimental results . . . . .	17
<b>6</b>	<b>Model 2: Online method for training a recurrent deep GSN</b>	<b>18</b>

6.1	Using real data sequence to train deep GSN . . . . .	18
6.2	Algorithm . . . . .	18
6.3	Extending the walkback procedure to sequenced inputs . . . . .	18
6.4	Analysis . . . . .	18
<b>7</b>	<b>Model 3: Hybrid method for training a recurrent deep GSN</b>	<b>19</b>
7.1	Generalizing the EM model . . . . .	19
7.2	Algorithm . . . . .	19
7.3	Experimental results . . . . .	19
<b>8</b>	<b>Discussion</b>	<b>20</b>
<b>9</b>	<b>Conclusion</b>	<b>21</b>

# List of Tables

# List of Figures

3.1	An example deep architecture. . . . .	5
3.2	Principal component analysis. . . . .	6
3.3	A Restricted Boltzmann machine. . . . .	8
3.4	Stacked RBM. . . . .	9
3.5	Stacked auto-encoder. . . . .	10
4.1	GSN computational graph. . . . .	15
4.2	Unrolled GSN Markov chain. . . . .	15

# Chapter 1

## Introduction

Deep learning research has grown in popularity due to its ability to form useful feature representations of highly complex input data. Useful representations are those that disentangle the factors of variation of input data, preserving the information that is ultimately useful for the given machine learning task. Deep learning frameworks (especially deep convolutional neural networks [?]) have had recent successes for supervised learning of representations for many tasks, creating breakthroughs for both speech and object recognition [?, ?].

Unsupervised learning of representations, however, has had slower progress. These models, mostly Restricted Boltzmann Machines (RBM) [?], auto-encoders [?], and sparse-coding variants [?], suffer from the difficulty of marginalizing across an intractable number of configurations of random variables (observed, latent, or both). Each plausible configuration of latent and observed variables would be a mode in the distribution of interest  $P(X, H)$  or  $P(X)$  directly, and current methods of inference or sampling are forced to make strong assumptions about these distributions. Recent advances on the generative view of denoising auto-encoders and generative stochastic networks [?] have alleviated this difficulty by only having to learn a local Markov chain transition operator through simple backpropagation, which is often unimodal (instead of having to parameterize the data distribution directly, which is multi-modal). This approach has opened up



unsupervised learning of deep representations for many useful tasks, including sequence prediction. Unsupervised sequence prediction and labeling remains an important problem for AI, as many types of input data naturally form sequences and the vast majority is unlabeled, such as language, video, etc.

This thesis will cover four main topics:

- Chapter 3 provides an overview of deep architectures - a background on representation learning from probabilistic and direct encoding viewpoints. Recent work on generative viewpoints will be discussed as well, showing how denoising auto-encoders can solve the multi-modal problem via learning a Markov chain transition operator.
- Chapter 4 introduces Generative Stochastic Networks - recent work generalizing the denoising auto-encoder framework into GSNs will be explained, and how this can be extended to sequence prediction tasks.
- Chapters 5, 6, and 7 describe models using GSNs to learn complex sequential input data, providing experimental results on synthetic MNIST data.
- Chapter 8 is a discussion of all three models and why they are able to use deep representations to learn sequence data.

# Chapter 2

## Related Work

Due to the success of deep architectures on highly complex input data, applying deep architectures to sequence prediction tasks has been studied extensively in literature. RBM variants have been the most popular for applying deep learning models to sequential data.

*Temporal RBM (TRBM)* [?] is a framework

*Recurrent Temporal RBMs (RTRBMs)* [?] are an extension of the TRBM.

*Temporal Convolution Machines (TCMs)* [?] also build from TRBMs.

*RNN-RBM* [?] is similar to the RTRBM.

*Sequential Deep Belief Networks (SDBNs)* [?, ?]

*Recursive Neural Networks (RNNs)* [?] are a slightly different framework used for sequence labeling in parsing sentences or image scenes.

Past work has also compared a deep architecture, Sentence-level Likelihood Neural Nets (SLNN), with traditional Conditional Random Fields (CRF) for sequence labeling tasks of Named Entity Recognition and Syntactic chunking [?]. Wang et al. found that non-linear deep architectures are more effective in low dimensional continuous input spaces, but are not more effective in high-dimensional discrete input spaces compared to linear CRFs. They also confirm that distributional representations can be used to achieve better generalization.

# Chapter 3

## Background: Deep Architectures

Current machine learning algorithms' performance depend heavily on the particular features of the data chosen as inputs. For example, document classification (such as marking emails as spam or not) can be performed by breaking down the input document into bag-of-words or n-grams as features. Choosing the correct feature representation of input data, or feature engineering, is a way that people can bring prior knowledge of a domain to increase an algorithm's computational performance and accuracy. To move towards general artificial intelligence, algorithms need to be less dependent on this feature engineering and better learn to identify the explanatory factors of input data on their own [3].

### 3.1 Representation Learning

Deep learning frameworks (also known as deep architectures) move in this direction by capturing a good representation of input data by using compositions of non-linear transformations. A good representation can be defined as one that disentangles underlying factors of variation for input data [2]. Deep learning frameworks can find useful abstract representations of data across many domains: it has had great commercial success powering most of Google and Microsoft's current speech recognition, image classification, natural language processing, object recognition, etc. Facebook is also planning on using deep learning to understand

its users<sup>1</sup>. Deep learning has been so impactful in industry that MIT Technology Review named it as a top-10 breakthrough technology of 2013<sup>2</sup>.

The central idea to building a deep architecture is to learn a hierarchy of features one level at a time where the input to one computational level is the output of the previous level for an arbitrary number of levels. Otherwise, shallow representations (such as regression or support vector machines) go directly from input data to output classification.

One good analogue for deep architectures is neurons in the brain (a motivation for artificial neural networks) - the output of a group of neurons is agglomerated as the input to more neurons to form a hierarchical layer structure. Each layer  $N$  is composed of  $h$  computational nodes that connect to each computational node in layer  $N + 1$ .




Figure 3.1: An example deep architecture.

## 3.2 Interpretations of Deep Architectures

There are two main ways to interpret the computation performed by these layered deep architectures:

- *Probabilistic graphical models* have nodes in each layer that are considered as latent random variables. In this case, you care about the probability

<sup>1</sup><http://www.technologyreview.com/news/519411/facebook-launches-advanced-ai-effort-to-find-meaning-in-your-posts/>

<sup>2</sup><http://www.technologyreview.com/featuredstory/513696/deep-learning/>

distribution of the input data  $x$  and the hidden latent random variables  $h$  that describe the input data in the joint distribution  $p(x, h)$ . These latent random variables describe a distribution over the observed data.

- *Direct encoding models* have nodes in each layer that are considered as computational units. This means each node  $h$  performs some computation (normally nonlinear like a sigmoidal function, hyperbolic tangent nonlinearity, or rectifier linear unit) given its inputs from the previous layer.

To get started, principal component analysis (PCA) is a simple feature extraction algorithm that can span both of these interpretations. PCA learns a linear transform  $h = f(x) = W^T x + b$  where  $W$  is a weight matrix for the inputs  $x$  and  $b$  is a bias term. The columns of the  $dx \times dh$  matrix  $W$  form an orthogonal basis for the  $dh$  orthogonal directions of greatest variance in the input training data  $x$ . The result is  $dh$  features that make representation layer  $h$  that are decorrelated.

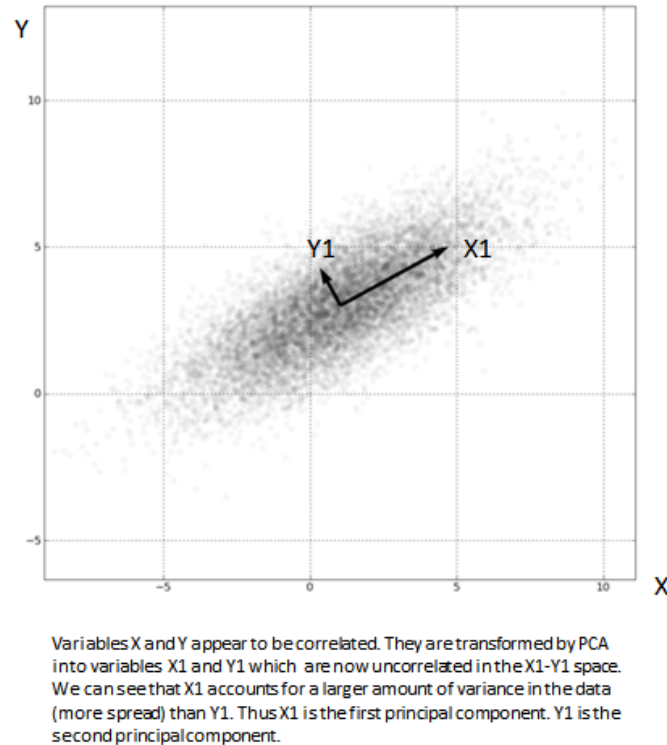


Figure 3.2: Principal component analysis.

From a probabilistic viewpoint, PCA is simply finding the principal eigenvectors of the covariance matrix of the data. This means that you are finding which features of the input data can explain away the most variance in the data[1]. From an encoding viewpoint, PCA is performing a linear computation over the input data to form a hidden representation  $h$  that has a lower dimensionality than the data.

Note that because PCA is a linear transformation of the input  $x$ , it cannot really be stacked in layers because the composition of linear operations is just another linear operation. There would be no abstraction benefit of multiple layers. To show these two methods of analysis, this section will examine stacking Restricted Boltzmann Machines (RBM) from a probability viewpoint and nonlinear auto-encoders from a direct encoding viewpoint.

### 3.2.1 Probabilistic models: restricted boltzmann machine (RBM)

A Boltzmann machine is a network of symmetrically-coupled binary random variables or units. This means that it is a fully-connected, undirected graph. This graph can be divided into two parts:

1. The visible binary units  $x$  that make up the input data and
2. The hidden or latent binary units  $h$  that explain away the dependencies between the visible units  $x$  through their mutual interactions.

Boltzmann machines describe this pattern of interaction through the distribution over the joint space  $[x, h]$  with the energy function:

$$\varepsilon_{\Theta}^{BM}(x, h) = -\frac{1}{2}x^T U x - \frac{1}{2}h^T V h - x^T W h - b^T x - d^T h$$

Where the model parameters  $\Theta$  are  $\{U, V, W, b, d\}$ .

Trying to evaluate conditional probabilities over this fully connected graph ends up being an intractable problem. For example, computing the conditional probability of hidden variable given the visibles,  $P(h_i|x)$ , requires marginalizing over all the other hidden variables. This would be evaluating a sum with  $2dh - 1$  terms.

However, we can restrict the graph from being fully connected to only containing the interactions between the visible units  $x$  and hidden units  $h$ .




Figure 3.3: A Restricted Boltzmann machine.

This gives us an RBM, which is a *bipartite* graph with the visible and hidden units forming distinct layers. Calculating the conditional distribution  $P(h_i|x)$  is readily tractable and now factorizes to:

$$P(h|x) = \prod_i P(h_i|x)$$

$$P(h_i = 1|x) = \text{sigmoid} \left( \sum_j W_{ji} x_j + d_i \right)$$

Very successful deep learning algorithms stack multiple RBMs together, where the hidden units  $h$  from the visible input data  $x$  become the new input data for another RBM for an arbitrary number of layers.

There are a few drawbacks to the probabilistic approach to deep architectures:




Figure 3.4: Stacked RBM.

1. The posterior distribution  $P(h_i|x)$  becomes incredibly complicated if the model has more than a few interconnected layers. We are forced to resort to sampling or approximate inference techniques to solve the distribution, which has computational and approximation error prices.
2. Calculating this distribution over latent variables still does not give a usable feature vector to train a final classifier to make this algorithm useful for AI tasks. For example, we calculate all of these hidden distributions that explain the variations over the handwriting digit recognition problem, but they do not give a final classification of a number. Actual feature values are normally derived from the distribution, taking the latent variable's expected value, which are then used as the input to a normal machine learning classifier, such as logistic regression.

### 3.2.2 Direct encoding models: auto-encoder

To get around the problem of deriving useful feature values, an auto-encoder is a non-probabilistic alternative approach to deep learning where the hidden units produce usable numeric feature values. An auto-encoder directly maps an input  $x$  to a hidden layer  $h$  through a parameterized closed-form equation called an encoder. Typically, this encoder function is a nonlinear transformation of the



input to  $h$  in the form:

$$f_{\Theta}(x) = s_f(b + Wx)$$

This resulting transformation is the feature-vector or representation computed from input  $x$ . Conversely, a decoder function is used to then map from this feature space  $h$  back to the input space, which results in a reconstruction  $x'$ . This decoder is also a parameterized closed-form equation that is a nonlinear undoing the encoding function:

$$g_{\Theta}(h) = s_g(d + W'h)$$

In both cases, the nonlinear function  $s$  is normally an element-wise sigmoid, hyperbolic tangent nonlinearity, or rectifier linear unit.

Thus, the goal of an auto-encoder is to minimize a loss function over the reconstruction error given the training data. Model parameters  $\Theta$  are  $\{W, b, W', d\}$ , with the weight matrix  $W$  most often having tied weights such that  $W' = W^T$ .

Stacking auto-encoders in layers is the same process as with RBMs.




Figure 3.5: Stacked auto-encoder.

One disadvantage of auto-encoders is that they can easily memorize the training data (i.e. find the model parameters that map every input seen to a perfect reconstruction with zero error) given enough hidden units  $h$ . To combat this problem, regularization is necessary, which gives rise to variants such as sparse auto-encoders, contractive auto-encoders, or denoising auto-encoders.

A practical advantage of auto-encoder variants is that they define a simple, tractable optimization objective that can be used to monitor progress.

### 3.3 Denoising Auto-encoders

Denoising auto-encoders [?, ?, ?] are a class of direct encoding models that use synthetic noise over the inputs through a corruption process during training to prevent overfitting and simply learning the identity function. Given a known corruption process  $C(\tilde{X}|X)$  to corrupt an observed variable  $X$ , the denoising auto-encoder learns the reverse conditional  $P(X|\tilde{X})$ . Combining this estimator with the known corruption process  $C$ , it can recover a consistent estimator of  $P(X)$  through a Markov chain that alternates sampling from  $C(\tilde{X}|X)$  and  $P(X|\tilde{X})$ . The basic algorithm is as follows:

---

**Algorithm 1:** Generalized Denoising Auto-encoder Training Algorithm

---

**Input:** training set  $D$  of examples  $X$ , a corruption process  $C(\tilde{X}|X)$ , and a conditional distribution  $P_{\Theta}(X|\tilde{X})$  to train.

**while** *training not converged* **do**

- sample training example  $X \sim D$ ;
- sample corrupted input  $\tilde{X} \sim C(\tilde{X}|X)$ ;
- use  $(X, \tilde{X})$  as an additional training example towards minimizing the expected value of  $-\log P_{\Theta}(X|\tilde{X})$ , e.g., by a gradient step with respect to  $\Theta$  in the encoding/decoding function;

**end**

---

The reconstruction distribution  $P(X|\tilde{X})$  is easier to learn than the true data distribution  $P(X)$  because  $P(X|\tilde{X})$  is often dominated by a single or few major modes, where the data distribution  $P(X)$  would be highly multimodal and complex. Recent works [?, ?] provide proofs that denoising auto-encoders with arbitrary variables (discrete, continuous, or both), an arbitrary corruption (Gaussian or other; not necessarily asymptotically small), and an arbitrary loss function (as long as it is viewed as a log-likelihood) estimate the score (derivative of the log-density with respect to the input) of the observed random variables.

Another key idea presented in Bengio et al. [?] is walkback training. The walkback process generates additional training examples through a pseudo-Gibbs sampling process from the current denoising auto-encoder Markov chain for a certain number of steps. These additional generated  $(X, \tilde{X})$  pairs from the model

decrease training time by actively correcting spurious modes (regions of the input data that have been insufficiently visited during training, which may therefore be incorrect in the learned reconstruction distribution). Both increasing the number of training iterations and increasing corruption noise alleviate spurious modes, but walkbacks are the most effective.

---

**Algorithm 2:** Walkback Training Algorithm for Denoising Auto-encoders

---

**Input:** A given training example  $X$ , a corruption process  $C(\tilde{X}|X)$ , and the current model’s reconstruction conditional distribution  $P_{\Theta}(X|\tilde{X})$ . It also has a hyper-parameter  $p$  that controls the number of generated samples.

**Output:** A list  $L$  of additional training examples  $\tilde{X}^*$ .

$X^* \leftarrow X, L \leftarrow [];$

Sample  $\tilde{X}^* \sim C(\tilde{X}|X^*);$

Sample  $u \sim \text{Uniform}(0, 1);$

**while**  $u < p$  **do**

Append  $\tilde{X}^*$  to  $L$ , so  $(X, \tilde{X}^*)$  will be an additional training example for the denoising auto-encoder.;

Sample  $X^* \sim P_{\Theta}(X|\tilde{X}^*);$

Sample  $\tilde{X}^* \sim C(\tilde{X}|X^*);$

Sample  $u \sim \text{Uniform}(0, 1);$

**end**

Append  $\tilde{X}^*$  to  $L;$

Return  $L;$

---

# Chapter 4

## General Methodology: Deep Generative Stochastic Networks (GSN)

Generative stochastic networks are a generalization of the denoising auto-encoder and help solve the problem of mixing between many modes as outlined in the Introduction. Each model presented in this thesis uses the GSN framework for learning a more useful abstraction of the input distribution  $P(X)$ .

### 4.1 Generalizing denoising auto-encoders

Denoising auto-encoders use a Markov chain to learn a reconstruction distribution  $P(X|\tilde{X})$  given a corruption process  $C(\tilde{X}|X)$  for some data  $X$ . Denoising auto-encoders have been shown as generative models [?], where the Markov chain can be iteratively sampled from:

$$\begin{aligned} X_t &\sim P_{\Theta}(X|\tilde{X}_{t-1}) \\ \tilde{X}_t &\sim C(\tilde{X}|X_t) \end{aligned}$$

As long as the learned distribution  $P_{\Theta_n}(X|\tilde{X})$  is a consistent estimator of the

true conditional distribution  $P(X|\tilde{X})$  and the Markov chain is ergodic, then as  $n \rightarrow \infty$ , the asymptotic distribution  $\pi_n(X)$  of the generated samples from the denoising auto-encoder converges to the data-generating distribution  $P(X)$  (proof provided in Bengio et al. [?]).

#### 4.1.1 Easing restrictive conditions on the denoising auto-encoder

A few restrictive conditions are necessary to guarantee ergodicity of the Markov chain - requiring  $C(\tilde{X}|X) > 0$  everywhere that  $P(X) > 0$ . Particularly, a large region  $V$  containing any possible  $X$  is defined such that the probability of moving between any two points in a single jump  $C(\tilde{X}|X)$  must be greater than 0. This restriction requires that  $P_{\Theta_n}(X|\tilde{X})$  has the ability to model every mode of  $P(X)$ , which is a problem this model was meant to avoid.

To ease this restriction, Bengio et al. [?] proves that using a  $C(\tilde{X}|X)$  that only makes small jumps allows  $P_{\Theta}(X|\tilde{X})$  to model a small part of the space  $V$  around each  $\tilde{X}$ . This weaker condition means that modeling the reconstruction distribution  $P(X|\tilde{X})$  would be easier since it would probably have fewer modes.

However, the jump size  $\sigma$  between points must still be large enough to guarantee that one can jump often enough between the major modes of  $P(X)$  to overcome the deserts of low probability:  $\sigma$  must be larger than half the largest distance of low probability between two nearby modes, such that  $V$  has at least a single connected component between modes. This presents a tradeoff between the difficulty of learning  $P_{\Theta}(X|\tilde{X})$  and the ease of mixing between modes separated by this low probability desert.

#### 4.1.2 Generalizing to GSN

While denoising auto-encoders can rely on  $X_t$  alone for the state of the Markov chain, GSNs introduce a latent variable  $H_t$  that acts as an additional state variable

in the Markov chain along with the visible  $X_t$  [?]:

$$H_{t+1} \sim P_{\Theta_1}(H|H_t, X_t)$$

$$X_{t+1} \sim P_{\Theta_2}(X|H_{t+1})$$

The resulting computational graph takes the form:




Figure 4.1: GSN computational graph.

The latent state variable  $H$  can be equivalently defined as  $H_{t+1} = f_{\Theta_1}(X_t, Z_t, H_t)$ , a learned function  $f$  with an independent noise source  $Z_t$  such that  $X_t$  cannot be reconstructed exactly from  $H_{t+1}$ . If  $X_t$  could be recovered from  $H_{t+1}$ , the reconstruction distribution would simply converge to the dirac at  $X$ . Denoising auto-encoders are therefore a special case of GSNs, where  $f$  is fixed instead of learned.

GSNs also use the notion of walkback to aid training. The resulting Markov chain of a GSN is inspired by Gibbs sampling, but with stochastic units at each layer that can be backpropagated [?].




Figure 4.2: Unrolled GSN Markov chain.

## 4.2 Extension to recurrent deep GSN

Mixing problem: depth to the rescue. MCMC sampling from the top layer model. Deeper nets visit more modes faster. Deeper representations lead to better abstractions lead to disentangling. More abstract/disentangled representations unfold manifolds and fill more space -  $\mathcal{I}$  can be exploited for better mixing between modes.

# Chapter 5

## Model 1: Expectation

## Maximization method for training a recurrent deep GSN

The first approach makes the assumption that GSN's learn complex feature transformations that represent the most information-bearing aspects of the input. Therefore, the recurrent step should be less complicated (through a single regression step). This model uses both artificial time-sequencing (for the GSN training to have BPTT), as well as real time-sequencing (for regression step).

### 5.1 Recurrent nature of deep GSNs

### 5.2 Algorithm

### 5.3 Experimental results



# Chapter 6

## Model 2: Online method for training a recurrent deep GSN

The second approach utilizes past network state to train the recurrent step at the same time as the GSN step.

### 6.1 Using real data sequence to train deep GSN

### 6.2 Algorithm

### 6.3 Extending the walkback procedure to sequenced inputs

### 6.4 Analysis

# Chapter 7

## Model 3: Hybrid method for training a recurrent deep GSN

Both the EM and online learning methods for training the recurrent deep GSN have their drawbacks - EM only models the sequence linearly in the latent variable  $H$  space while the online learning method requires untied weights between layers. This model, the hybrid approach, aims to combine the best of both worlds by creating an alternating structure inspired by convolutional neural networks with alternating convolutional and pooling layers [?].

### 7.1 Generalizing the EM model

### 7.2 Algorithm

### 7.3 Experimental results

# Chapter 8

## Discussion

## Chapter 9

## Conclusion

# Bibliography

- [1] Francis R. Bach and Michael I. Jordan. A probabilistic interpretation of canonical correlation analysis. Technical Report 688, University of California, Berkeley, April 2005.
- [2] Yoshua Bengio. Deep learning of representations: Looking forward. *CoRR*, abs/1305.0445, 2013.
- [3] Yoshua Bengio, Aaron C. Courville, and Pascal Vincent. Unsupervised feature learning and deep learning: A review and new perspectives. *CoRR*, abs/1206.5538, 2012.